

TSBC-1585 Report - Burn-tail matrix cells

TSBC-1585 Progress Report - 2026-07-30

Hypothesis

The three deferred burn-tail matrix cells can be closed before the Friday 14:30 reset by rerunning the Grok 4.5 Paperclip R1 cell, then closing the Hermes-current engineer cell, then running the lean quant cell only after the first two cells are decision-grade.

Method

1. Started the TSBC-1547 follow-up path for the Grok 4.5 Paperclip R1 cell:
 - Runner:
`/Users/glad0s/.paperclip/instances/default/projects/e212ce50-b524-408c-b3d4-0c6108d8c2e2/f71e8665-3f38-4920-b777-348ec85b9071/_default/work-products/TSBC-1141/TSBC-1262/run-paperclip-r1.py`
 - Environment: `TSBC_R1_REPS=3`
 - Issue override: `ISSUE_IDENTIFIER=TSBC-1585` , `PARENT_IDENTIFIER=TSBC-1547`
 - Report stem: `TSBC-1585-paperclip-r1`
2. Ran the Hermes-current engineer standard probe frame:
 - Harness: `/Users/glad0s/paperclip/benchmark/tsbc_task_probe.py`
 - Config:
`/Users/glad0s/.paperclip/instances/default/companies/e212ce50-b524-408c-b3d4-0c6108d8c2e2/work-products/TSBC-1543/TSBC-1548/config/config-grok-4.5-hermes-current.json`
 - Role: `engineer`
 - Reps: `3`
 - Agent file: `current`
 - Skills: `none`
3. Did not run lean quant because the Paperclip R1 cell has not reached a decision-grade terminal result.

Data

Grok 4.5 Paperclip R1

- Attempt run root: `/Users/glad0s/paperclip/benchmark/work-products/TSBC-1585/evidence/paperclip/paperclip-20260730-195417/`
- Sample root: `/Users/glad0s/paperclip/benchmark/work-products/TSBC-1585/evidence/paperclip/TSBC-1585-paperclip-r1-20260730T195417Z/`
- Custody note: the sample-home evidence copy is sanitized; `.hermes` auth and env files were removed from the work-product copy before packaging.
- Created fixture: `TSBC-1588 ( e7655201-4235-464f-a1ff-c3b23e8f2096 )` titled `[agentic-bench] Stage 1 – read, comment, done`
- Created heartbeat run: `d02cfd0d-b64d-4ff7-a3fa-d79725d27551`
- Current state at report time: heartbeat `queued` , fixture `todo` , clean bench agent `Bench-grok-4.5-hermes-clean` idle.
- Scored rows: `0`
- Ledger rows: `0`
- Decision-grade verdict: no. The runner reached `paperclip_lane.run_case()` and waited on the queued Paperclip heartbeat; no terminal fixture result or R1 record was produced.

Hermes-current Engineer

- Probe run: `probe-20260730-205955`
- Result root: `/Users/glad0s/paperclip/benchmark/results/probe-20260730-205955/`
- Evidence copy: `/Users/glad0s/paperclip/benchmark/work-products/TSBC-1585/evidence/engineer-current/probe-20260730-205955/`
- Samples: `36`
- OK count: `36/36`
- Mean quality: `0.973113425925926`
- Minimum quality: `0.5720833333333334`
- Mean output tokens: `1969.9444444444443`
- Ledger rows appended: `12`
- Served-model evidence from raw rows: requested `grok-4.5-hermes-current` ; reported served model `grok-4.5` ; adapter `hermes` .
- Decision-grade verdict: yes, sufficient for the engineer-current cell despite one lower-quality sample, because all 36 samples completed and the aggregate mean quality is high.

Lean Quant

- Not run.
- Gating reason: issue contract says to run lean quant only after cells 1 and 2 are decision-grade. Cell 1 is not decision-grade yet.

Verdict

INCONCLUSIVE.

The Hermes-current engineer cell is closed and banked to the ledger, but the Grok 4.5 Paperclip R1 cell is still unscored because its Paperclip heartbeat is queued. The lean quant cell remains intentionally deferred until the Paperclip R1 cell is decision-grade.

Dispatched Follow-up Key

`TSBC-1585-R1-QUEUE-CONTINUATION`: resume the Paperclip R1 follow-up by either allowing heartbeat `d02cfd0d-b64d-4ff7-a3fa-d79725d27551` / fixture `TSBC-1588` to start and reconstructing the sample record, or cancelling that stale queued fixture and rerunning `TSBC_R1_REPS=3` through `run-paperclip-r1.py` once the clean bench lane can execute immediately. Do not run lean quant until that R1 cell is decision-grade.