

TSBC-1432 Bench Ledger Recording Closeout

TSBC-1432 Bench Ledger Recording Closeout

Source issue: [TSBC-1432](#)

Report date: 2026-07-28

Brand pack: stack-lab

Hypothesis verdict: CONFIRMED

Dispatched follow-up key: none required

Hypothesis

The benchmark harness can keep the existing repetition enforcement and publication gates while making the ledger auditable: every completed run should write per-pass evidence with requested/served model identity, timestamps, pass/fail, and failure reasons; aggregate rows should remain queryable but become derivable from pass rows; historical rows should be migrated with explicit `unknown` markers rather than guessed values.

Method

I treated the reopened comment as a rescope: reps were already enforced, so the implementation target became ledger recording and migration.

Implementation changes:

- Added pass start/finish timestamps and served-model verification fields to `benchmark/bench.py` run records.
- Added AGY same-run self-report probing in `benchmark/adapters.py`; pin/self-report mismatch returns `failureReason: served_model_mismatch` before task execution.
- Updated `benchmark/scoring.py` so served-model mismatches are never scored.
- Expanded `benchmark/ledger.py` with `model_eval_pass` rows, aggregate derivation from pass rows, `passes` query support, idempotent expanded-schema backfill, failure-reason classification, and low-success suppression preservation.
- Updated `benchmark/README.md`, `benchmark/TSBC-CHARTER.md`, and canonical TSKB0047 benchmark fairness protocol.

Live migration:

- Ran `python3 ledger.py backfill-expanded-schema` on `/Users/glad0s/paperclip/benchmark/ledger/results.jsonl`.
- Recovered pass rows only where `results/<run_id>/runs.json` existed.
- Marked unrecoverable historical fields as `unknown`.
- Marked every `gemini-pro` row `served_model_verified: false` under TSBC-1439.
- Preserved the success-rate floor backfill for the historical `agy-claude-opus-4.6 / cto` row.

Data

Live ledger after migration:

Metric	Value
Total ledger rows	7,872
<code>model_eval</code> aggregate rows	1,295
<code>model_eval_pass</code> rows	4,852
v2 <code>model_eval</code> rows	1,295
<code>model_eval</code> /pass rows with <code>reps</code> , <code>requested_model</code> , <code>served_model</code>	6,147
Failed non-skipped pass rows	37
Failed pass rows without <code>failure_reason</code>	0
Low-success aggregate rows	2
Low-success rows still carrying quality	0
Low-success rows with <code>suppressed_reason</code>	2
<code>gemini-pro</code> rows	241
<code>gemini-pro</code> rows with <code>served_model_verified: false</code>	241

Historical repetition handling:

- Aggregate rows with explicit recovered `reps` : 3.
- Aggregate rows with `reps: "unknown"` : 1,292.
- Pass rows with explicit recovered `reps` : 90.
- Pass rows with `reps: "unknown"` : 4,762.

The suppressed historical row now has:

- `quality: null`
- `qPer1kOut: null`

- `successRate: 0.0833`
- `suppressed_reason: "successRate 0.0833 below floor 0.8000"`
- `decision_band: "failed"`
- `pass_records_recovered: 12`
- `reps: "unknown" with observed_passes_per_task: 1`

Verification

Passed:

- `python3 -m unittest test_harness_gates.py -> 8 tests.`
- `python3 -m py_compile bench.py ledger.py adapters.py scoring.py report.py benchlib.py .`
- `python3 bench.py all --roles cto --models grok-4.3 --max-tasks-per-role 1 --dry-run -> prints reps : 3 (decision floor 3) .`
- `python3 - <<'PY' ... ledger.backfill_expanded_schema() ... PY -> { "changed": 0, "pass_rows_added": 0, "backup": null } .`
- `python3 ledger.py query cto agy-claude-opus-4.6 --days 120 --min 1 -> nDecisionGradeResults: 0 , suppressedResults: 1 , passRows: 12 .`
- `python3 ledger.py passes cto agy-claude-opus-4.6 --run run-20260725-202032 --days 120 -> 12 pass rows.`

Not run:

- Full repo typecheck/build/test. The change is scoped to the Python benchmark harness and the live benchmark ledger.

Verdict

CONFIRMED . The revised requirement is implemented: future completed benchmark runs write pass-level rows, queryable aggregate rows carry `reps` /requested/served model fields, historical rows are migrated with honest `unknown` markers, the known low-success rows are suppressed, and AGY pin/self-report mismatches are recorded as failures that are never scored.

Evidence

- Evidence root: `work-products/TSBC-1432/evidence/`
- Ledger counts: `work-products/TSBC-1432/evidence/ledger-counts.json`
- Idempotency proof: `work-products/TSBC-1432/evidence/backfill-idempotency.json`

- Low-success query: `work-products/TSBC-1432/evidence/query-cto-agy-opus-suppressed.txt`
- Historical pass rows: `work-products/TSBC-1432/evidence/passes-cto-agy-opus-run-20260725-202032.json`
- Explicit-reps pass rows: `work-products/TSBC-1432/evidence/passes-cv-review-grok-4-20-run-20260728-165618.json`
- Failure reason counts: `work-products/TSBC-1432/evidence/failure-reason-counts.json`
- Unit tests: `work-products/TSBC-1432/evidence/unittest-test-harness-gates.txt`
- Python compile: `work-products/TSBC-1432/evidence/py-compile-touched.txt`
- Dry-run proof: `work-products/TSBC-1432/evidence/bench-dry-run-reps.txt`
- Code/docs diff stat: `work-products/TSBC-1432/evidence/git-diff-stat.txt`
- Final mutating ledger backup: `benchmark/ledger/results.jsonl.bak-tsbc-1432-expanded-ledger-20260728164030`

TSKB

Canonical TSKB delta: `/Users/glad0s/TSKB/KB/TSKB0047 [TSBC] - Lane Placement and Skill-Pack Adoption Tracker - v0.1 - 07-06/benchmark-fairness-protocol.md`.

The update records the reusable rule that decision-grade benchmark evidence requires pass-level ledger rows, honest `unknown` markers for unrecoverable historical fields, and AGY same-run served-model self-report handling.

Render Command

```
brandsuite pdf --brand stack-lab -- --input work-products/TSBC-1432/TSBC-1432-repo
```