

TSBC-1571 Retired-Alias Provenance Closeout

TSBC-1571 Retired-Alias Provenance Closeout

Source issue: [TSBC-1571](#)

Report date: 2026-07-30

PDF filename: TSBC-1571-report.pdf

Brand pack: stack-lab

Hypothesis verdict: CONFIRMED

Dispatched follow-up key: none required

Hypothesis

Post-retirement xAI fast-alias benchmark rows can preserve the originally requested label without double-counting Grok 4.3, and the benchmark harness can prevent new retired-alias pins from silently passing served-model verification.

Method

The implementation used the operator-verified May 15, 2026 xAI retirement as the provenance boundary. The public xAI migration note maps the retired Grok 4 fast aliases to `grok-4.3` and maps `grok-code-fast-1` to `grok-build-0.1`; the local Hermes `xai-oauth` model catalog also excludes the retired aliases and exposes `grok-4.3` / `grok-build-0.1`.

Implementation steps:

1. Added `benchmark/model_provenance.py` as the single retired-alias map and ledger/report normalization helper.
2. Annotated the live benchmark ledger in place, preserving original model labels while adding correction fields.
3. Updated `benchmark/ledger.py`, `benchmark/report.py`, and `benchmark/costreport.py` so aggregate/ranking views fold corrected fast aliases into `grok-4.3`.
4. Updated `benchmark/adapters.py` so retired alias pins fail loudly with `failureReason: retired_model_alias`; `_served_model_matches_pin` no longer accepts token-overlap matches for retired aliases.

5. Removed retired aliases from active benchmark rosters, Paperclip-lane agent mappings, pricing keys, and helper-script defaults, while keeping the denylist/provenance/test references.
6. Updated the canonical model-watch KB and the live 07:30 model-watch routine so future config, lane, and benchmark-matrix changes check for the retired aliases before use.

Data

Ledger annotation result:

Field	Value
Rows annotated	2,003
<code>model_original = grok-4-fast</code>	864
<code>model_original = grok-4.1-fast</code>	1,139
Rows with <code>served_model_corrected = grok-4.3</code>	2,003
Effective <code>grok-4.3</code> rows after fold	2,594
Effective <code>grok-4-fast</code> rows after fold	0
Effective <code>grok-4.1-fast</code> rows after fold	0

Ledger backup:

- `benchmark/ledger/results.jsonl.bak-tsbc-1571-retired-alias-20260730185435`

Matcher diff:

- `benchmark/adapters.py` : rejects retired alias pins before live calls and rejects self-reported retired aliases after AGY probing.
- `benchmark/model_provenance.py` : central map for retired aliases, cutoff date, annotation fields, and failure payloads.
- Failure payload includes `failureReason: retired_model_alias`, `servedModelVerified: false`, and `servedModelMismatch: true`.
- `_served_model_matches_pin` returns false for retired pin aliases even when an API echoes the requested alias or overlaps tokens with the served model.

Aggregate/ranking fold:

- `benchmark/ledger.py` : `read_all`, `query`, and `passes` return effective model IDs for reporting without destructively rewriting history.

- `benchmark/report.py` : aggregate cells normalize rows through `model_effective / served_model_corrected` .
- `benchmark/costreport.py` : model normalization uses the same effective ID.

Active config hygiene:

Surface	Retired aliases remaining as runnable pins
<code>models / models_catalog ids</code>	0
<code>model_arg values</code>	0
Paperclip agentic lane agent keys	0
Pricing keys	0

Expected remaining references are limited to the denylist, provenance map, model-watch runbook text, and tests.

Verification

Passed:

- `python3 -m unittest test_harness_gates.py` from `benchmark/` -> 12 tests.
- `python3 -m py_compile benchmark/adapters.py benchmark/ledger.py benchmark/report.py benchmark/costreport.py benchmark/model_provenance.py benchmark/cascade.py benchmark/per_task_compare.py benchmark/team_bench.py benchmark/smoke_cheap.py benchmark/run_disposition_shadow.py benchmark/tsbc_713_coverage.py` .
- `python3 -m json.tool benchmark/config.json >/dev/null` .
- `python3 bench.py all --dry-run` -> default roster includes `grok-4.3` and `grok-4.20` , and excludes the retired fast aliases.
- Config hygiene assertion -> retired aliases absent from active model ids, model args, Paperclip agent keys, and pricing keys.
- Ledger fold assertion -> `grok-4.3: 2594` , `grok-4-fast: 0` , `grok-4.1-fast: 0` , annotated read rows `2003` .
- `python3 ledger.py query content grok-4-fast --days 60` -> query folds the requested retired alias to `grok-4.3` and records `requestedModelQuery` .

Not run:

- Full repo typecheck/build/test. The change is scoped to the Python benchmark harness, benchmark configs, the live benchmark ledger, and the model-watch routine/KB.

Verdict

CONFIRMED . Historical post-cutoff fast-alias rows now preserve request provenance while aggregating under the served Grok 4.3 identity, new retired alias pins fail served-model verification loudly, and model-watch now blocks future lane/config/bench pins for the retired aliases.

Evidence

- Work-product root: `work-products/TSBC-1571/`
- Report source: `work-products/TSBC-1571/TSBC-1571-report.md`
- PDF render: `work-products/TSBC-1571/TSBC-1571-report.pdf`
- Ledger backup: `benchmark/ledger/results.jsonl.bak-tsbc-1571-retired-alias-20260730185435`
- Source doc: `https://docs.x.ai/developers/migration/may-15-retirement`
- Hermes catalog evidence: `/Users/glad0s/.hermes/provider_models_cache.json`
- Canonical KB reference: `/Users/glad0s/TSKB/KB/TSKB0056 [ALL] - Model-Watch & New-Model Intake Runbook - v1.0 - 07-11.md`
- KB index reference: `/Users/glad0s/TSKB/KB/INDEX.md`
- Live routine updated: `94bec39d-c427-4016-a282-de2eb5e83151`

Gate Checks

- Directive completeness: no **DEVIATION** ; ledger, matcher, and model-watch requirements were all implemented.
- Payload verification: exact retired-alias matcher and ledger-fold behavior were exercised in focused tests and direct ledger assertions.
- Routine semantics: only the 07:30 model-watch checklist text changed; cadence and schedule ownership were not changed.
- Artifact custody: source and PDF render are preserved under `work-products/TSBC-1571/` ; the PDF is uploaded before issue close.

Never-again gates checked: G2, G3, G7, G8, G9, G10. Evidence: this report and the attached Paperclip PDF/source artifacts. Unsatisfied gates: none.

Render Command

```
brandsuite pdf --brand stack-lab -- --input work-products/TSBC-1571/TSBC-1571-repo
```